

5.1)

```
for (i=0; i<8; i++)
  for (j=0; j<8000; j++)
    A[i][j] = B[i][0] + A[i][j];
```

5.1.1 How many 32-bit ints in 16 byte cache block?

4

5.1.2 Which variables exhibit temporal locality?

j, i

5.1.3 Spatial locality

C codes gets stored rows first


	0	1	2
0			
1			

 3x2 so we have spatial locality when accessing several items on the same row.

```
for i=1:8
  for j=1:8000
    A(i,j) = B(i,0) + A(j,i);
  end
end
```

5.1.4 How many 16-byte cache blocks are needed to store all 32-bit elems referenced?
 $(8 \times 8000 + 8000 + 8000 \times 8 - 8 \times 8) \cdot \frac{1}{4} = 33984$

5.1.5 Temporal
i, j



5.1.6 Spatial
A(j,i), B(j,0)


5.5)

512 KiB working set with address stream: 0 2 4 6 8 10 12 ...

5.5.1 Assume 64 KiB direct mapped cache with 32-byte block. What is the MR?
 How is it sensitive to the set size or cache size?

$$\frac{64}{32} = 2K \text{ blocks}$$

1 bytes per elem
 32 bytes per block



First access => Miss, read 16 elems to cache (1, 3, 5, ... garbage)
 17th access => — 11 —

Miss every 16th access. Only cold misses due to streams not having locality.

5.5.2 MR for 16-byte, 64 and 128-byte block size

16 byte: $\frac{1}{8}$
 64 byte: $\frac{1}{32}$
 128 byte: $\frac{1}{64}$

What kind of locality is exploited?
 Spatial. We read sequentially.

5.5.3 Assume a two-entry stream buffer.

First read would give a miss, then no misses.

Cache block size (B) can affect both miss rate and miss latency. Assuming CPI=1 and an avg of 1.35 references per instruction and the optimal block size given the following MR for block sizes:

B	8	16	32
MR	4%	3%	2%

5.5.4 What is the optimal block size for miss latency of $20 \times B$ cycles?

B=8: $0.04(8 \cdot 20)$
 B=16: $0.03(16 \cdot 20)$ 8 wins!
 ...

5.5.5 $24+B$ cycles. 24 units of time to fetch the first elem and the consecutive elems takes 1 unit

B=8: $0.04(24+8)$
 B=16: $0.03(24+16)$
 ...

5.5.6 For const miss latency?
 Only the miss rate matters!

5.11

5.11 As described in Section 5.7, virtual memory uses a page table to track the mapping of virtual addresses to physical addresses. This exercise shows how this table must be updated as addresses are accessed. The following data constitutes a stream of virtual addresses as seen on a system. Assume 4 KiB pages, a 4-entry fully associative TLB, and true LRU replacement. If pages must be brought in from disk, increment the next largest page number.

4669, 2227, 13916, 34587, 48870, 12608, 49225

TLB

Valid	Tag	Physical Page Number
1	11	12
1	7	4
1	3	6
0	4	9

Page table

Valid	Physical Page or in Disk
1	5
0	Disk
0	Disk
1	6
1	9
1	11
0	Disk
1	4
0	Disk
0	Disk
1	3
1	12

5.11.1 Show the final stage of the system given the above info. Also show if it's a hit in TLB, Page table or page fault.

4669: 0001 0010 0011 1101 12 bits
 Tag Byte offset in page 4 KiB pages $\Rightarrow 4 \cdot 2^{10} \Rightarrow 12$ bits to encode
 Tag 1 is not in TLB \Rightarrow TLB miss
 (Tag = Virtual Page)

5.11.2 Use 16 KiB instead of pages. Advantages / Disadvantages?

- + Fewer misses since you load larger chunks.
- Larger misspenalty!